

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import cross_val_score, KFold
from sklearn.preprocessing import StandardScaler
```

```
In [2]: df=pd.read_csv('AAPL.csv')
#Basic Data information
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10468 entries, 0 to 10467
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    Date        10468 non-null   object
1    Open         10468 non-null   float64
2    High         10468 non-null   float64
3    Low          10468 non-null   float64
4    Close        10468 non-null   float64
5    Adj Close    10468 non-null   float64
6    Volume       10468 non-null   int64
dtypes: float64(5), int64(1), object(1)
memory usage: 572.6+ KB
```

```
In [3]: df
```

```
Out[3]:
```

	Date	Open	High	Low	Close	Adj Close	Volume
0	1980-12-12	0.128348	0.128906	0.128348	0.128348	0.100178	469033600
1	1980-12-15	0.122210	0.122210	0.121652	0.121652	0.094952	175884800
2	1980-12-16	0.113281	0.113281	0.112723	0.112723	0.087983	105728000
3	1980-12-17	0.115513	0.116071	0.115513	0.115513	0.090160	86441600
4	1980-12-18	0.118862	0.119420	0.118862	0.118862	0.092774	73449600
...
10463	2022-06-13	132.869995	135.199997	131.440002	131.880005	131.880005	122207100
10464	2022-06-14	133.130005	133.889999	131.479996	132.759995	132.759995	84784300
10465	2022-06-15	134.289993	137.339996	132.160004	135.429993	135.429993	91533000
10466	2022-06-16	132.080002	132.389999	129.039993	130.059998	130.059998	108123900
10467	2022-06-17	130.070007	133.080002	129.809998	131.559998	131.559998	134118500

10468 rows × 7 columns

```
In [4]: # Display summary statistics
print(df.describe())
```

	Open	High	Low	Close	Adj Close
count	10468.000000	10468.000000	10468.000000	10468.000000	10468.000000
mean	14.757987	14.921491	14.594484	14.763533	14.130431
std	31.914174	32.289158	31.543959	31.929489	31.637275
min	0.049665	0.049665	0.049107	0.049107	0.038329
25%	0.283482	0.289286	0.276786	0.283482	0.235462
50%	0.474107	0.482768	0.465960	0.475446	0.392373
75%	14.953303	15.057143	14.692589	14.901964	12.835269
max	182.630005	182.940002	179.119995	182.009995	181.511703

	Volume
count	1.046800e+04
mean	3.308489e+08
std	3.388418e+08
min	0.000000e+00
25%	1.237768e+08
50%	2.181592e+08
75%	4.105794e+08
max	7.421641e+09

In [5]:

```
# Check for missing values
print(df.isnull().sum())
```


```
Date          0
Open          0
High          0
Low           0
Close         0
Adj Close     0
Volume        0
dtype: int64
```

In [6]:

```
# Time Series analysis of the data and the target variable
# Convert 'Date' column to datetime format
df['Date'] = pd.to_datetime(df['Date'])

# Set 'Date' as the index
df.set_index('Date', inplace=True)

# Plot time series data
plt.figure(figsize=(12, 6))
plt.plot(df['Close'], label='Apple Stock Close Price')
plt.title('Apple Stock Close Price Over Time')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.show()
```



In [7]:

```
df.corr()['Close']
```


Out[7]:

```
Open          0.999850
High          0.999924
Low           0.999928
Close         1.000000
Adj Close     0.999671
Volume       -0.196411
Name: Close, dtype: float64
```

```
In [8]: # Calculate correlation matrix

correlation_matrix = df.corr()

# Plot a heatmap of the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



In [9]: *# Scatter plots wit all h variables*

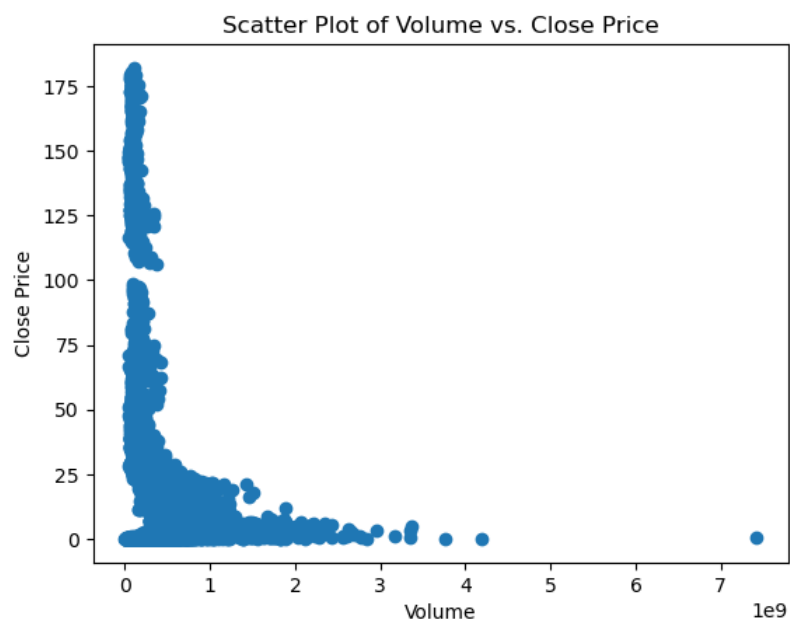
```
plt.scatter(df['Volume'], df['Close'])
plt.title('Scatter Plot of Volume vs. Close Price')
plt.xlabel('Volume')
plt.ylabel('Close Price')
plt.show()


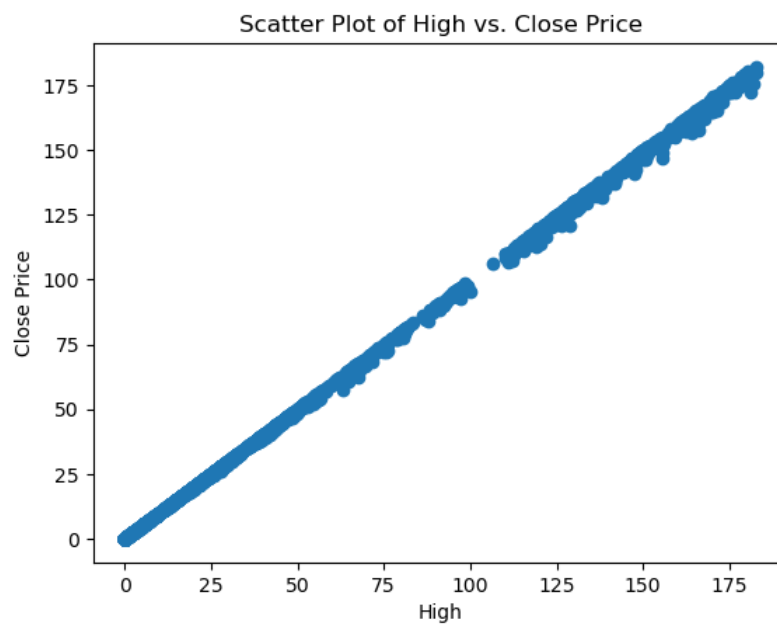
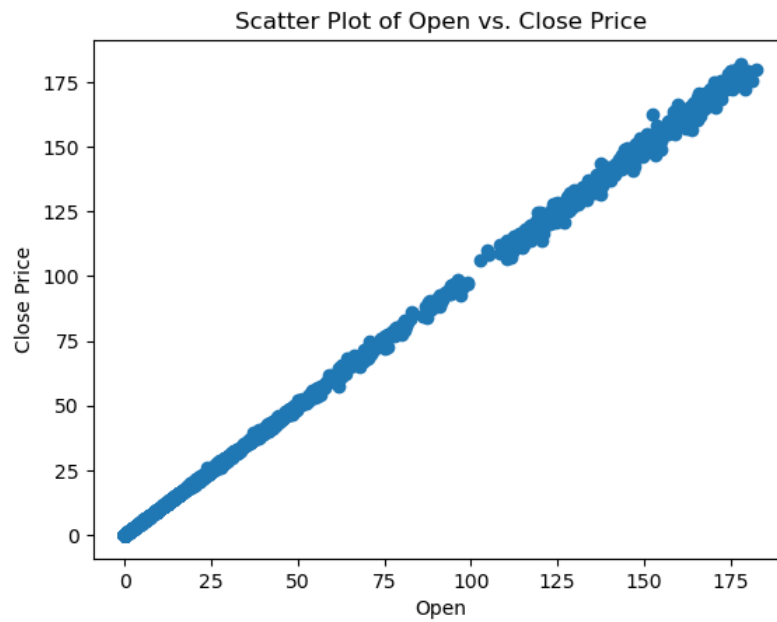
plt.scatter(df['Open'], df['Close'])
plt.title('Scatter Plot of Open vs. Close Price')
plt.xlabel('Open')
plt.ylabel('Close Price')
plt.show()

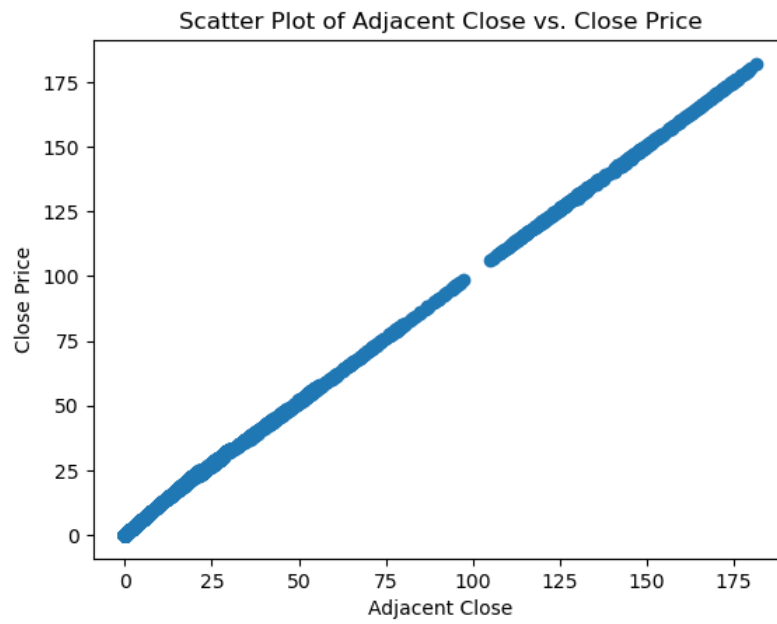
plt.scatter(df['High'], df['Close'])
plt.title('Scatter Plot of High vs. Close Price')
plt.xlabel('High')
plt.ylabel('Close Price')
plt.show()

plt.scatter(df['Low'], df['Close'])
plt.title('Scatter Plot of Low vs. Close Price')
plt.xlabel('Low')
plt.ylabel('Close Price')
plt.show()

plt.scatter(df['Adj Close'], df['Close'])
plt.title('Scatter Plot of Adjacent Close vs. Close Price')
plt.xlabel('Adjacent Close')
plt.ylabel('Close Price')
plt.show()
```







```

In [10]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import KFold, cross_val_score, train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error, r2_score

# Load the dataset (assuming 'AAPL.csv' contains the necessary columns)
df = pd.read_csv('AAPL.csv')

# Assume 'Close' is the column representing stock prices
data = df[['Open', 'High', 'Low', 'Adj Close']] # Exclude 'Close' from predictors
target = df['Close'].values.reshape(-1, 1)

# Normalize the data
scaler_data = MinMaxScaler(feature_range=(0, 1))
scaler_target = MinMaxScaler(feature_range=(0, 1))

data_normalized = scaler_data.fit_transform(data)
target_normalized = scaler_target.fit_transform(target)

# Function to prepare the data for linear regression
def create_dataset(data, target, look_back=1):
    X, y = [], []
    for i in range(len(data) - look_back):
        X.append(data[i:(i + look_back), :])
        y.append(target[i + look_back, 0])
    return np.array(X), np.array(y)

# Set the look-back period (number of time steps to look back)
look_back = 20

# Create the dataset
X, y = create_dataset(data_normalized, target_normalized, look_back)

# Set the number of folds for cross-validation
k_folds = 5

# Create a k-fold cross-validation object
kf = KFold(n_splits=k_folds, shuffle=True, random_state=42)

# Initialize the linear regression model
model = LinearRegression()

# Initialize lists to store MSE and R-squared scores
mse_scores = []
r2_scores = []

# Perform cross-validation and obtain scores
for train_index, test_index in kf.split(X):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]

    # Train the model on the training set
    model.fit(X_train.reshape(-1, look_back * data.shape[1]), y_train)

    # Make predictions on the test set
    y_pred = model.predict(X_test.reshape(-1, look_back * data.shape[1]))

    # Denormalize the predictions and actual values
    y_pred_denormalized = scaler_target.inverse_transform(y_pred.reshape(-1, 1))
    y_test_denormalized = scaler_target.inverse_transform(y_test.reshape(-1, 1))

    # Calculate Mean Squared Error (MSE) on the test set
    mse_test = mean_squared_error(y_test_denormalized, y_pred_denormalized)

    # Calculate R-squared on the test set
    r2_test = r2_score(y_test_denormalized, y_pred_denormalized)

    mse_scores.append(mse_test)
    r2_scores.append(r2_test)

# Print the mean and standard deviation of the MSE and R-squared scores
print(f'Mean MSE: {np.mean(mse_scores)}')
print(f'Standard Deviation MSE: {np.std(mse_scores)}')
print(f'Mean R-squared: {np.mean(r2_scores)}')
print(f'Standard Deviation R-squared: {np.std(r2_scores)}')

# No need to train the model again on the entire dataset after cross-validation

# Split the data into training and testing sets

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=False)

# Make predictions on the test set
y_pred = model.predict(X_test.reshape(-1, look_back * data.shape[1]))

# Denormalize the predictions and actual values
y_pred_denormalized = scaler_target.inverse_transform(y_pred.reshape(-1, 1))
y_test_denormalized = scaler_target.inverse_transform(y_test.reshape(-1, 1))

# Calculate Mean Squared Error (MSE) on the test set
mse_test = mean_squared_error(y_test_denormalized, y_pred_denormalized)

# Calculate R-squared on the test set
r2_test = r2_score(y_test_denormalized, y_pred_denormalized)

print(f'Mean Squared Error (MSE) on Test Set: {mse_test}')
print(f'R-squared on Test Set: {r2_test}')

# Plot actual vs. predicted values on the test set
# Plot actual vs. predicted values on the test set

df['Date'] = pd.to_datetime(df['Date'])

# Plot for actual values and the predictions.
plt.figure(figsize=(12, 6))
plt.plot(df['Date'][-len(y_test_denormalized):], y_test_denormalized, label='Actual Closing Price', color='blue')
plt.plot(df['Date'][-len(y_test_denormalized):], y_pred_denormalized, label='Predicted Closing Price', color='orange')
plt.title('Apple Stock Close Price Over Time')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.show()

```

Mean MSE: 0.5365784710577322
 Standard Deviation MSE: 0.07546522768126811
 Mean R-squared: 0.999475381709283
 Standard Deviation R-squared: 6.340824985322405e-05
 Mean Squared Error (MSE) on Test Set: 2.1335148126186203
 R-squared on Test Set: 0.9989716080512189

